

COMP90015
Distributed Systems
Project 1 report

Group: PerpetualMotionSquad

Members: Malcolm Karutz – mkarutz@student.unimelb.edu.au

Alexander Zable – azable@student.unimelb.edu.au

Brandon Syiem – bsyiem@student.unimelb.edu.au

Ishesh Gambhir – igambhir@student.unimelb.edu.au

1. Introduction

The aim of this project was to develop a multi-server message broadcast middleware by implementing the protocol defined in the project specification. The client API allows applications to register and login to the system with a unique username – and broadcast arbitrary JSON objects, called activities, to all other clients connected to the network. The system features load-balancing for incoming client connections, synchronisation strategies to ensure consistent client registration across the multi-server network, and authentication features to prevent malicious servers from joining the network. Testing and debugging the consistency of the distributed system posed the greatest challenge in the project. The asynchronous, message-oriented protocol allowed a natural approach to writing and unit-testing handlers for each type of command, but validating load-balancing and user registration required extensive integration testing. Overall, we have achieved a relatively faithful implementation of the specification, however further integration testing is certainly warranted.

2. Server Failure Model

Handling graceful server crashes could involve extending the JSON protocol to include a `SERVER_QUIT` message, with which a receiving server could knowingly terminate the existing connection. Ultimately a `SERVER_QUIT` message could contain the address/port of its “parent” server, making reconnection to the network trivial (see Fig. 1). The main edge case

here is if the root server quits. In this case the root server could potentially designate one of its child servers as the new “root”, with which the rest of its children would have to connect to.

If a server can quit without warning, this complicates things. The TCP/IP protocol has a timeout period which would have to be reached before a remote server could be considered crashed. Then, we could potentially exploit prior `SERVER_ANNOUNCE` data to establish a connection to another server in the network. Perhaps even more reliably, we could have a master server (or a couple for redundancy) which simply maintains the server topology. Any server can query it if it needs to reconnect.

Up until now we have assumed a fail-stop model. If we assume that a server may restart at the same address and port number, we could potentially keep polling it by attempting to reconnect continually. However, given that this is a *may*, we are probably best off connecting to another server in the network as described above, at least in the meantime. We could continue to check for the reinstatement of the crashed server if maintaining the original tree structure is of importance.

Potential server failure may compromise our user registration model. If a server quits, the tree is at least temporarily split into two sections. Hence a




Figure 1. Server tree whereby a server fails gracefully, enabling its clients to connect to its parent.

LOCK_REQUEST will only propagate within the respective section it is initiated in. Therefore in the case of eventual reinstatement of the entire tree, the local storages will potentially be out of sync. Ideally servers should synchronize their states in the event of server reinstatement/connection.

3. Concurrency Issues

The protocol we implemented is far from flawless under certain edge cases. Let's consider two clients trying to register at the same time, with the same username and different secrets. One connects with a server and sends a registration request. The server it is connected to sends out a lock request that will eventually propagate to every other server after some delay. At the same time, on the "other side" of the tree of servers, the other client attempts to register with the same username but different secret. This server also propagates a lock request. In due course the lock requests sent will cause a clash, and a lock denied will be instigated for one of the username/secret combinations. However, until this happens there is a short time whereby one client could theoretically *log in* to another server which is only aware of *their* LOCK_REQUEST as yet. Thus, one user may be logged in as the other user who successfully registered (see Fig. 2).

A solution would be to keep a separate list of pending users on each server. The idea would be that a lock request adds a user to this list instead (and users on this list cannot log in). When the server that originally got the register request gets a lock allowed from every other server, it could then send a final broadcast (i.e. FINALIZE_REG), which would cause all servers to move the user from the pending list to the registered list.

Another concern is that load balancing as currently implemented may not be able to handle many login attempts because of the delay between SERVER_ANNOUNCE commands. A surge of client connections to a particular server could result in most of the surge being redirected to another single server, which will then redirect most of the remaining surge over and again. Some clients could experience many redirects before connecting. The most simple solution would be to increase the redirect tolerance (from 2) so that only when a server has *substantially* less connections do clients get redirected to it.

Concurrency issues such as these may be exacerbated when we consider the pathological case of a series of servers which form a "line". A broadcast could take a very long time to propagate to every server. This would be unacceptable in most applications, and ultimately care needs to be taken when setting up the server tree.

4. Scalability

If we model the server tree as a graph, we can then model a broadcast by counting each edge of the graph




Figure 2. Two clients register the same username, one with the "blue" secret and the other with the "green" secret. The bold outlined servers are vulnerable to being logged into during the registration process with the respectively coloured secret. Note how a third client can theoretically log in with the blue secret, even though eventually the green secret successfully wins the "lock battle", and registers successfully.

once (each edge "equivalent" to a single message sent). It is known that a tree with n nodes will have $n - 1$ edges. Therefore, $n - 1$ messages are required for a single broadcast, making a broadcast an $O(n)$ "message complexity" operation, where n is the number of servers.

Currently there are a few cases where all servers will broadcast to *each other*, such as the SERVER_ANNOUNCE every 5 seconds. Given there are n servers, this means there will be $n * (n - 1)$ messages sent throughout the network every 5 seconds; or in other words, $O(n^2)$ messages per second. Furthermore, if we consider that a lock request propagates to every server, and a LOCK_ALLOW or LOCK_DENIED reply broadcast is given in response, this would also be quadratic message complexity.

Evidently this would result in an unsustainable number of messages as the server network grows. A possible solution would involve utilizing a specialised master




Figure 3. Each server in the network periodically updates the "master server" with its load. In this example a client wishing to connect would be directed to S1.

server using a dynamic allocation policy (with status updates from the servers – see fig 2.). However, static policies such as round-robin have been shown to be only negligibly inferior to their dynamic counter-parts [1]. Therefore network overhead from server announcements can be eliminated completely, assuming user registrations are also handled differently (as below). A natural downside of any "master server" type approach such as this, is that a master server could potentially bottleneck the entire network, and would be a single point of failure. However in theory, a master server could be extended to become an abstract "service" which is a scalable distributed system itself.

User registrations could be handled similarly by utilising a specialised user authentication server/service. However, a more conservative solution would be to better exploit the tree structure of the network. We propose a revised registration protocol that acts recursively with servers only relying on lock information from their immediate neighbours, rather than every server in the network. Consider a server that receives a REGISTER command as the root of the tree; if it does not already know about the requested username, it will send a single LOCK_REQUEST to each of its immediate children servers and waits to receive LOCK_ALLOWED messages from them. Servers that receive a LOCK_REQUEST command, if they do not already know about the requested username, will recursively send LOCK_REQUEST messages to each of their immediate children and wait to receive LOCK_ALLOWED messages from each of them. When all the expected LOCK_ALLOWED messages are received, a server can respond to its parent with a LOCK_ALLOWED. At this point, the server requires confirmation from its parent before registering the user. We therefore propose that when the originating server receives all of the expected LOCK_ALLOWED or commands, or a single LOCK_DENIED command from any of its immediate children, it will broadcast a REGISTER_CONFIRM or REGISTER_ABORT

command, respectively. With this revised protocol, registering a user will require a maximum of $O(n)$ LOCK_REQUEST messages, $O(n)$ LOCK_ALLOWED messages, and $O(n)$ REGISTER_CONFIRM messages.

5. References

- [1] Y. Zhang, H. Kameda and S. L. Hung, "Comparison of dynamic and static load-balancing strategies in heterogeneous distributed systems," in *IEE Proceedings - Computers and Digital Techniques*, vol. 144, no. 2, pp. 100-106, Mar 1997. doi: 10.1049/ip-cdt:19970951